

Lab 04 – Building a Mapping Data Flow

This lab assumes you've created Lab 3 and have imported the following tables into your ADLS Gen 2 instance:

- SalesLT.Product
- SalesLT.ProductCategory
- SalesLT.ProductModel

Note: If your data factory instance is fairly old, it might not have the preview Mapping Data Flow activity. For the purpose of this lab – create a new data factory following Lab 01.

We're going to take a fairly basic warehousing example – take three tables and combine them into a basic “dimension” table.

We will use the lake copy of the Adventureworks database we used in previous examples – let's assume we're pulling the “name” category from the SalesLT.Product, SalesLT.ProductCategory and SalesLT.ProductModel tables into a single “dimension”-style table.

Basically, we're duplicating this query:

```
select P.ProductID,  
       P.Name ProductName,  
       PC.Name ProductCategory,  
       PM.Name ProductModel  
from SalesLT.Product P  
     inner join SalesLT.ProductCategory PC on P.ProductCategoryID =  
PC.ProductCategoryID  
     inner join SalesLT.ProductModel PM on P.ProductModelID = PM.ProductModelID
```

LAB 04.A Create Datasets

1. Data Flows are a little awkward as they're still in preview – they can't connect to an ADLS Gen 2 lake using a service principal just yet. So create a new Linked Service that uses the “Account Key” style of connection to your lake, just for this lab



Lab 04 – Building a Mapping Data Flow

Connect via integration runtime *

AutoResolveIntegrationRuntime

Authentication method

Account key

Account selection method

☐ From Azure subscription ☒ Enter manually

URL *

https://mdwlake.dfs.core.windows.net

Storage account key Azure Key Vault

Storage account key *

.....

- Just as we have before, we need to create three datasets pointing to files within our lake. Create a Data Lake Store Gen 2 dataset, for a CSV file and select the Product tables we imported earlier:

Name

DS_ADLS_SalesLTProduct

Linked service *

ADLS_MDWLake

[Edit Connection](#)

File path

lakeroot / RAW/Public/Adventure / SalesLT.Product.csv [Browse](#)

First row as header ☒

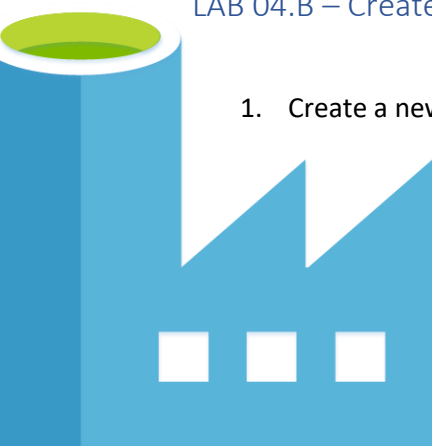
Import schema

☒ From connection/store ☐ From local file ☐ None


- Repeat the above steps for ProductCategory and ProductModel

LAB 04.B – Create a Data Flow

- Create a new Data Flow



Lab 04 – Building a Mapping Data Flow



- You'll be guided to add a source – click on the empty square and it will create a “source” stream component. Name this “Product” and select the DS_ADLS_SalesLTProduct dataset

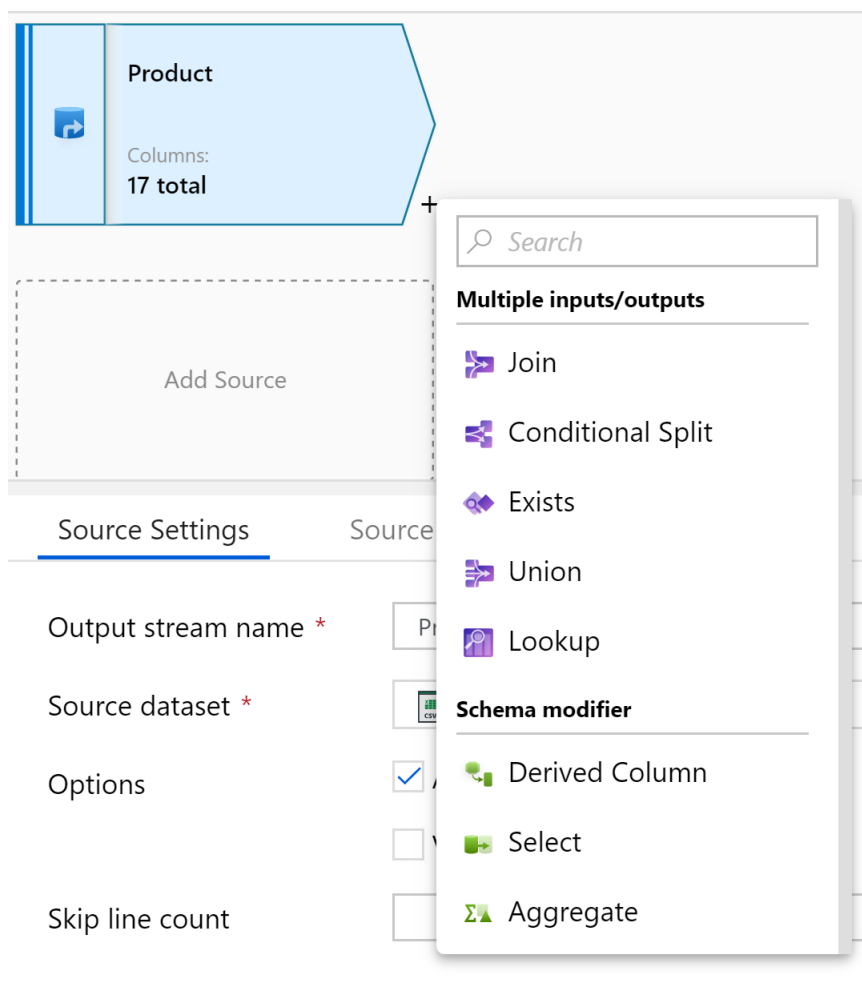
The screenshot displays the Azure Data Factory interface. At the top, a source stream component named 'Product' is shown with a blue icon and the text 'Columns: 17 total'. Below it is a dashed box labeled 'Add Source'. The bottom section shows the 'Source Settings' tab with the following configuration:

Source Settings	Source Options	Projection	Optimize	Inspect	Data Preview
Output stream name *	Product		Documentation		
Source dataset *	DS_ADLS_SalesLTProduct		Edit + New		
Options	<input checked="" type="checkbox"/> Allow schema drift ⓘ <input type="checkbox"/> Validate schema ⓘ				

- Click the small + icon next to the source stream to add a transformation:



Lab 04 – Building a Mapping Data Flow



4. We want to get rid of excess columns, so let's use a "Select" transformation. Give the transformation step a name:

Lab 04 – Building a Mapping Data Flow

✓ Validate

Product's column	Name as
abc ProductID	ProductID
abc Name	Name
abc ProductNumber	ProductNumber

- We want to rename “Name” to “Product” and remove all columns except for ProductID, ProductCategoryID, ProductModelID and our renamed “Product” attribute.

Select Settings Optimize Inspect Data Preview Description


Output stream name * SelectProductFields Documentation

Incoming stream * Product

Input columns * Auto Mapping Re-map Delete 4 items: 13 column(s) from the inputs left unmapped

Product's column	Name as
abc ProductID	ProductID
abc Name	Product
abc ProductCategoryID	ProductCategoryID
abc ProductModelID	ProductModelID

- We now want to click on the “Add Source” button to add a source for our next table:



- Configure this new source to use our ProductCategory dataset:

Lab 04 – Building a Mapping Data Flow

ProductCat
Columns: ProductID, ProductCategoryID, ProductCategory, ProductModelID, ProductModel

Source Settings Source Options Projection Optimize Inspect Data Preview

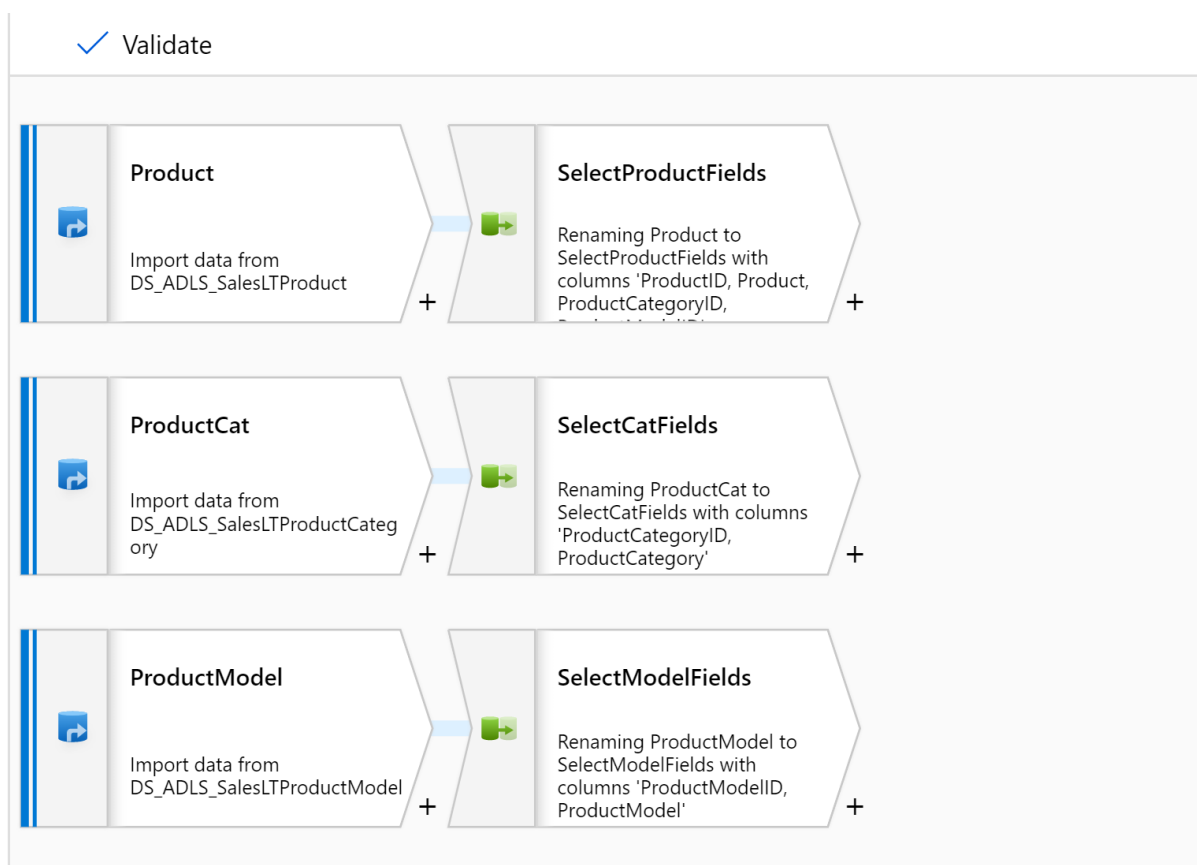
Output stream name * [Documentation](#)

Source dataset * DS_ADLS_SalesLTProductCategory Edit New

- Add Select Transform to rename “Name” to “ProductCategory” and remove all columns except for the ProductID

<input type="checkbox"/> ProductCat's column	Name as
<input type="checkbox"/> abc ProductCategoryID	<input type="text" value="ProductCategoryID"/>
<input type="checkbox"/> abc Name	<input type="text" value="ProductCategory"/>

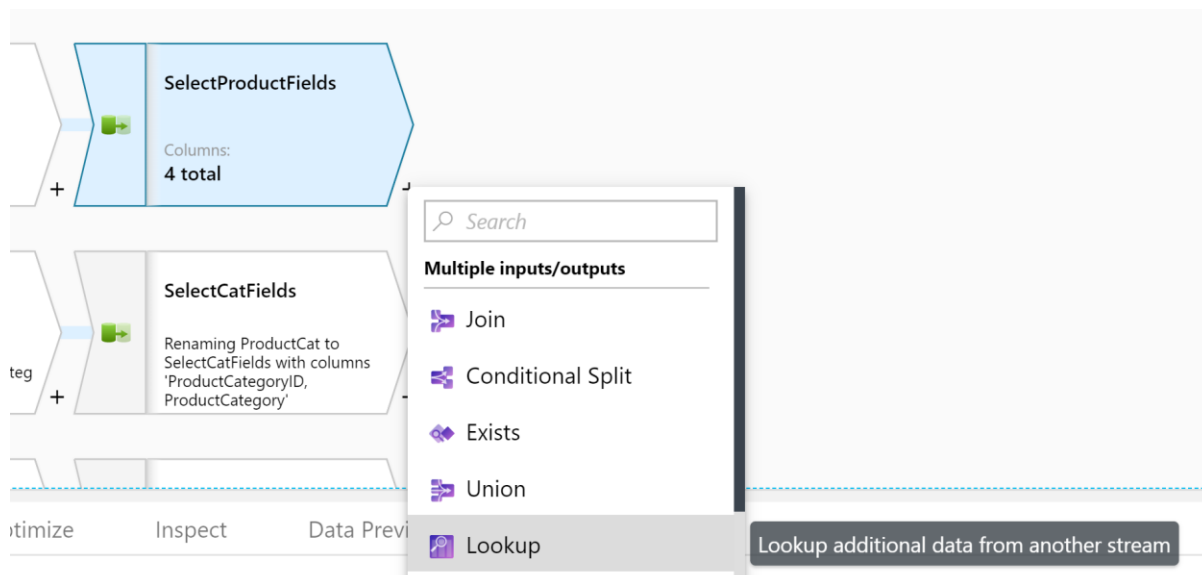
- Do the same for the ProductModel, again stripping out columns aside from the ProductModelID and the name (renamed as ProductModel)



Lab 04 – Building a Mapping Data Flow

We now have our source data, nicely trimmed of excess fields that we can combine into a single dataset.

10. Add a new transformation after the “select” on the main Product stream and choose the “Lookup” type – this is where we will lookup information from the other streams



11. Select the “SelectCatFields” stream as the reference stream and configure it to use the ProductCategoryID to perform the reference join, like so:

The 'Lookup Settings' tab is active, showing the following configuration:

- Output stream name *: [Documentation](#)
- Primary stream *: [SelectProductFields](#)
- Lookup stream *: [SelectCatFields](#)
- Lookup conditions *:

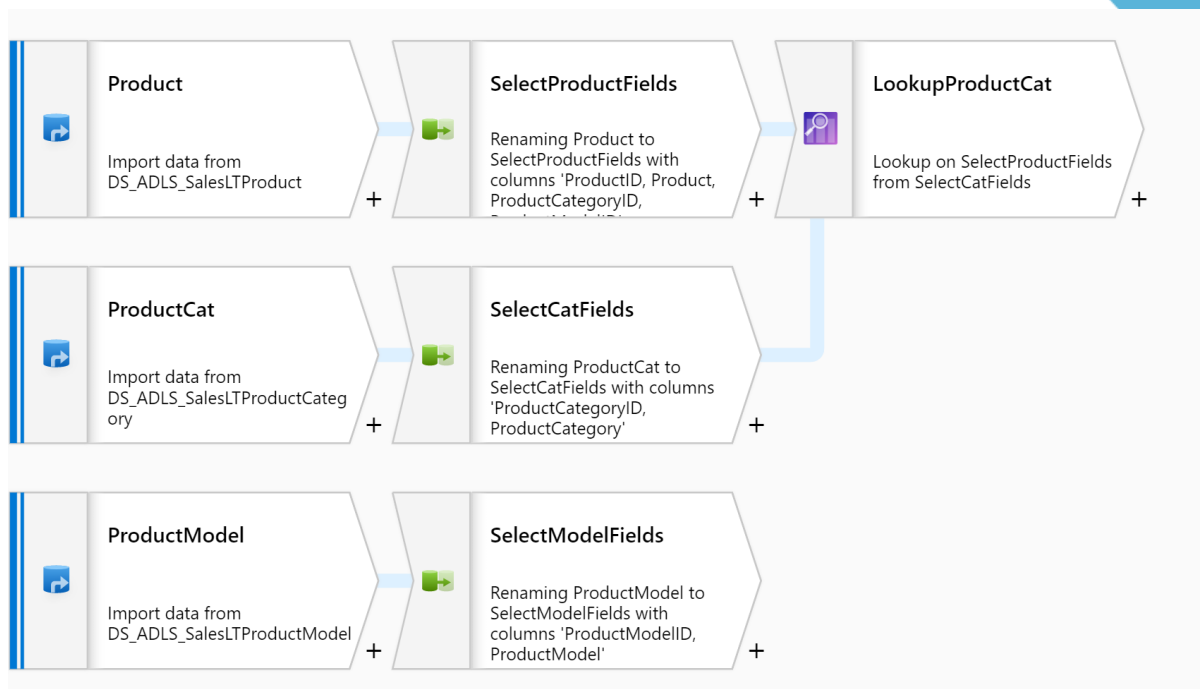
Left: SelectProductFields's column	=	Right: SelectCatFields's column
<input type="text" value="abc ProductCategoryID"/>		<input type="text" value="abc ProductCategoryID"/>

Be careful here – you can actually select to use the output of any of the transformations, including the original source ones, before we had stripped columns and renamed them.

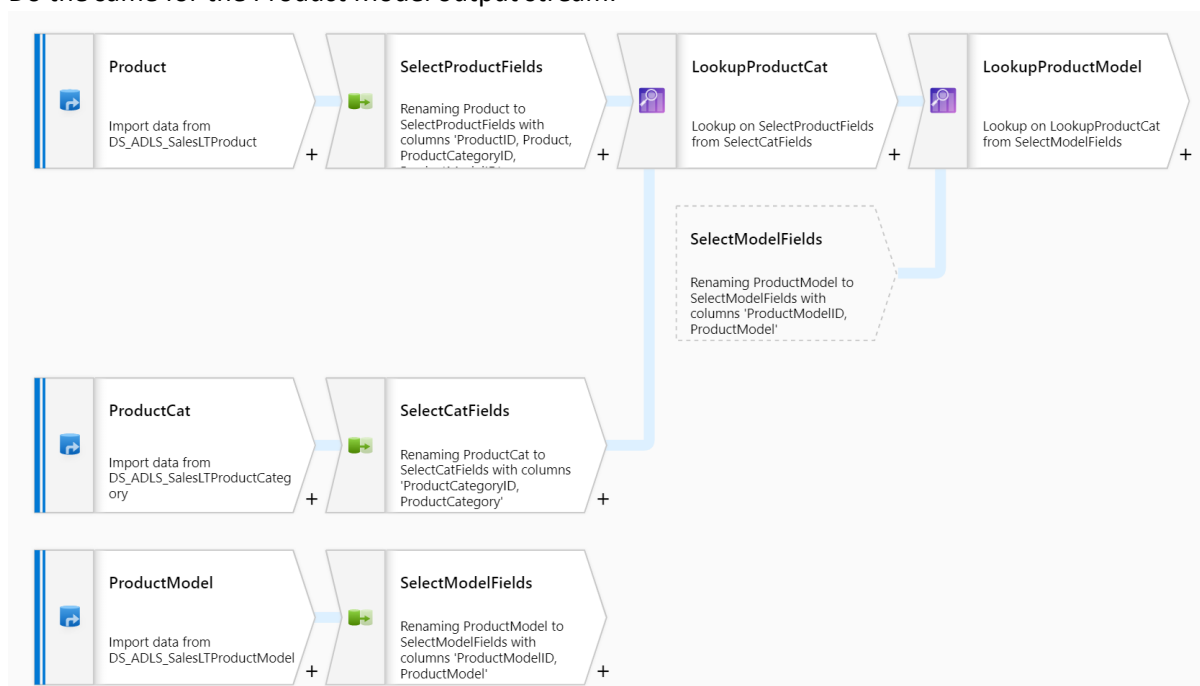
When your lookup is configured, it will automatically update your diagram reflect the relationship.



Lab 04 – Building a Mapping Data Flow



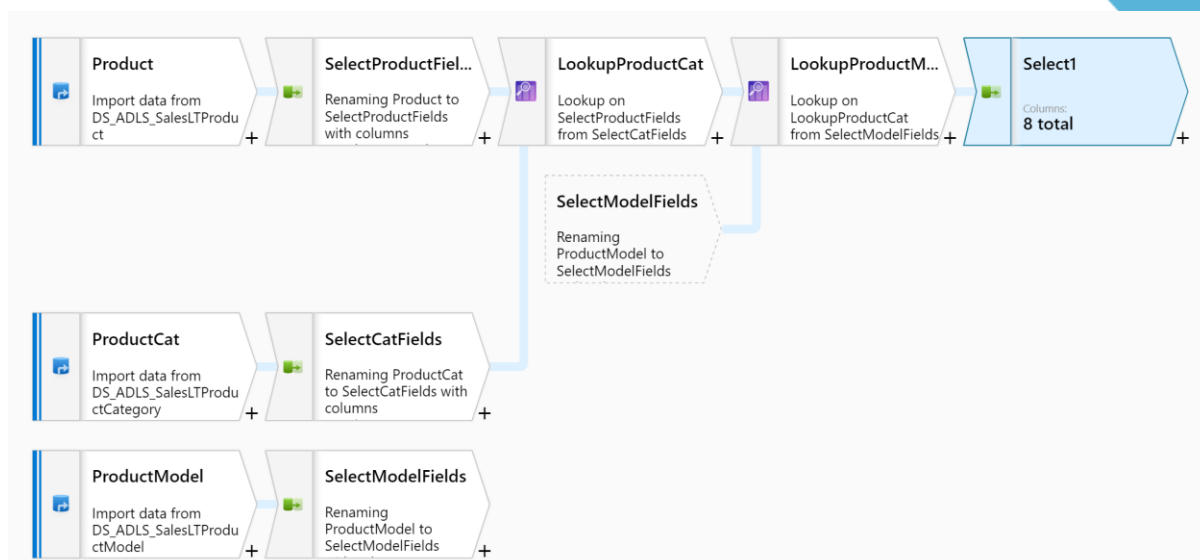
Do the same for the Product Model output stream:



You may notice, we didn't select the fields to be added at any point, just those that are the join constraints. By default, it will bring the whole table into the aggregation, including duplicates of the keys.

12. Let's add a final select transformation to get rid of those duplicate keys:

Lab 04 – Building a Mapping Data Flow



The duplicate columns will be highlighted and use a SELECT syntax to denote which stream they originally came from:

<input type="checkbox"/>	abc SelectProductFields@ProductCategoryID	→	! ProductCategoryID
<input type="checkbox"/>	abc SelectProductFields@ProductModelID	→	! ProductModelID
<input type="checkbox"/>	abc SelectCatFields@ProductCategoryID	→	! ProductCategoryID

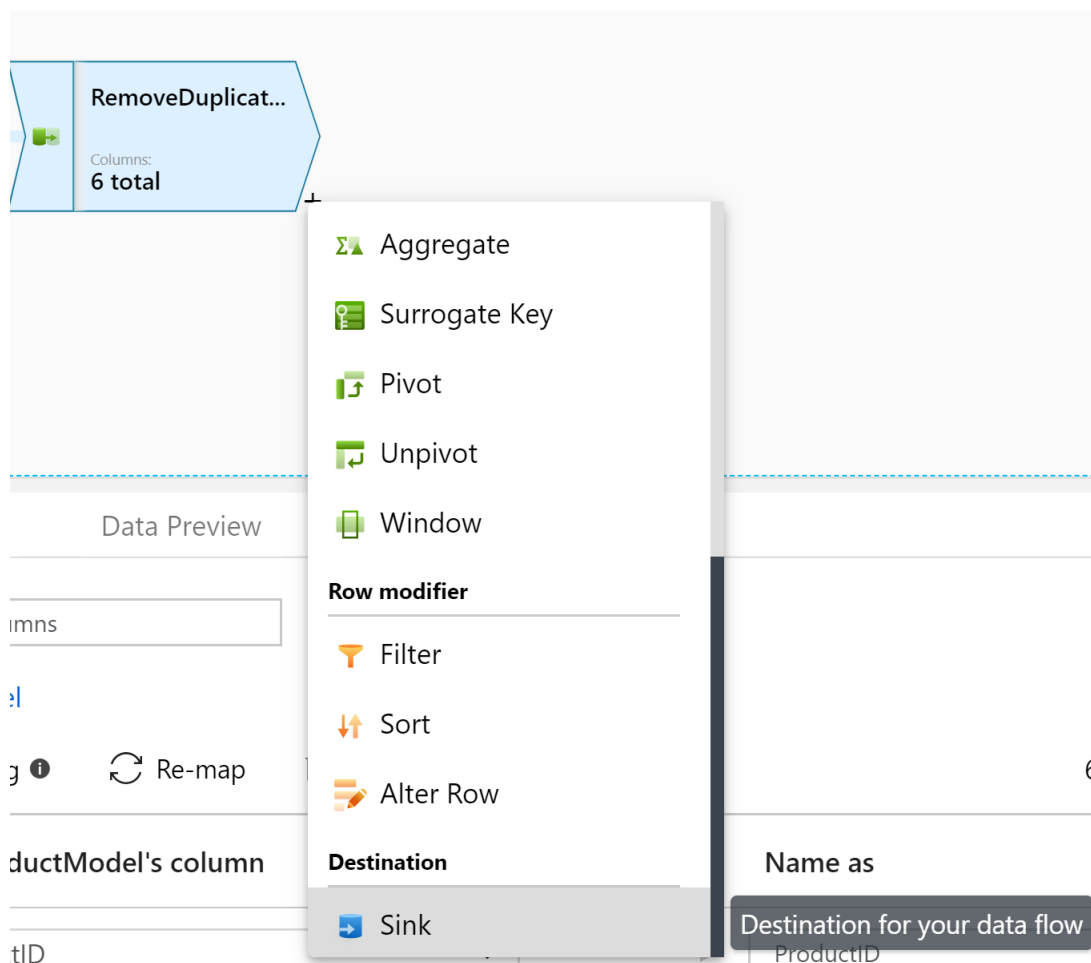
Delete the columns that are duplicates sourced from our lookup tables, and we should see something like:

<input type="checkbox"/>	abc ProductID	→	ProductID
<input type="checkbox"/>	abc Product	→	Product
<input type="checkbox"/>	abc SelectProductFields@ProductCategoryID	→	ProductCategoryID
<input type="checkbox"/>	abc SelectProductFields@ProductModelID	→	ProductModelID
<input type="checkbox"/>	abc ProductCategory	→	ProductCategory
<input type="checkbox"/>	abc ProductModel	→	ProductModel

- Finally, we want to write our data back to our lake in it's new form – we do this by adding a “Sink” transformation at the point where the data is in the right state:



Lab 04 – Building a Mapping Data Flow



We can choose an existing dataset if we had set one up in advance, or we can create one now, using the schema of our stream as reference

14. Name your sink transformation, then click the “New” button to create a new dataset to hold our data

Sink	Settings	Mapping	Optimize	Inspect	Data Preview
Output stream name *	WriteToADLS			Documentation	
Incoming stream *	RemoveDuplicateColumns				
Sink dataset *	Select...			+ New	
Options	<input checked="" type="checkbox"/> Allow schema drift ? <input type="checkbox"/> Validate schema ?				

Lab 04 – Building a Mapping Data Flow

15. Select Data Lake Store Gen 2 as the destination type, then select the file format.

For most scenarios, we would write this as a parquet file for performance – but for ease of testing our transformation, let's leave it as a CSV

You'll notice that several other big data formats are shown (ORC, Avro etc) but they are not yet enabled

16. Give the file a name and configure where it should be created within your lake. By convention, I've separated mine from the RAW data into a CURATED data layer:

Name
DS_ADLS_DimProduct

Linked service *
ADLS_MDWLake

[Edit Connection](#)

File path
lakeroot / CURATED/Warehouse / File [Browse](#)

First row as header ☒

Import schema
☒ From connection/store
 ☐ From local file
 ☐ None

Finish creating your dataset and navigate back to your data flow, you'll see it now updated with your sink:

Sink	Settings	Mapping	Optimize	Inspect	Data Preview
Output stream name *	WriteToADLS				Documentation
Incoming stream *	RemoveDuplicateColumns				
Sink dataset *	DS_ADLS_DimProduct				Edit + New
Skip line count					
Options	<input checked="" type="checkbox"/> Allow schema drift <input type="checkbox"/> Validate schema				

And that's it! That's our working Data Flow, ready to go!


Lab 04 – Building a Mapping Data Flow

Lab 04.C – Create a pipeline for your Data Flow

So... we've created a data flow and hooked it up to source data, destinations etc... but we've got nothing to actually run the data flow logic. That's where we need a pipeline.

1. Create a new pipeline and drag on a "Data Flow" activity, found in the "Move & Transform" menu:

▲ Move & Transform



2. Unlike other transformations, this will immediately open up a config window, where you'll need to select the name of the data flow you created earlier:

Adding Data Flow



- ☒ Use existing Data Flow
 ☐ Create new Data Flow

Existing Data Flow *

DF_Create_DimProduct ▼

If you look at the settings for your new activity, there isn't much to do. By default (and as the only option in preview), the data flow will work on an internal ADF Databricks cluster and perform its own sizing

+ - 🔒 100% 🔍 🖱️ 🔄 📄

General **Settings** User Properties

Data Flow * DF_Create_DimProduct ▼ Edit + New




Run on * AutoResolveIntegrationRuntime ▼ ⓘ

▶ PolyBase ⓘ

Lab 04 – Building a Mapping Data Flow




In future, it is expected that we'll be able to tweak the performance by changing the size/scale of the spark cluster our data flow is running on.

- Now we can test our new creation! Hit the publish button and trigger your pipeline to see how it goes


<input type="checkbox"/>	Pipeline Name ▾	Actions	Run Start ↕	Duration	Triggered By	Status	Parameters
	PL_CreateDimProductDF	 	06/27/2019, 1:49:49 PM	00:02:06	Manual trigger	 In Progress	

It'll probably take a couple of minutes before it does anything – that's because it's provisioning the spark cluster, which has an overhead. They're looking to reduce this, but for now bear in mind that these flows are generally meant for fairly large data processing tasks.





Eventually, we'll see this turn green as the transformation succeeds:


<input type="checkbox"/>	Pipeline Name ▾	Actions	Run Start ↕	Duration	Triggered By	Status
	PL_CreateDimProductDF	 	06/27/2019, 1:49:49 PM	00:06:19	Manual trigger	 Succeeded

What's good to learn is the debug/audit information available

- Click on the  icon to view the activity-level results for your pipeline

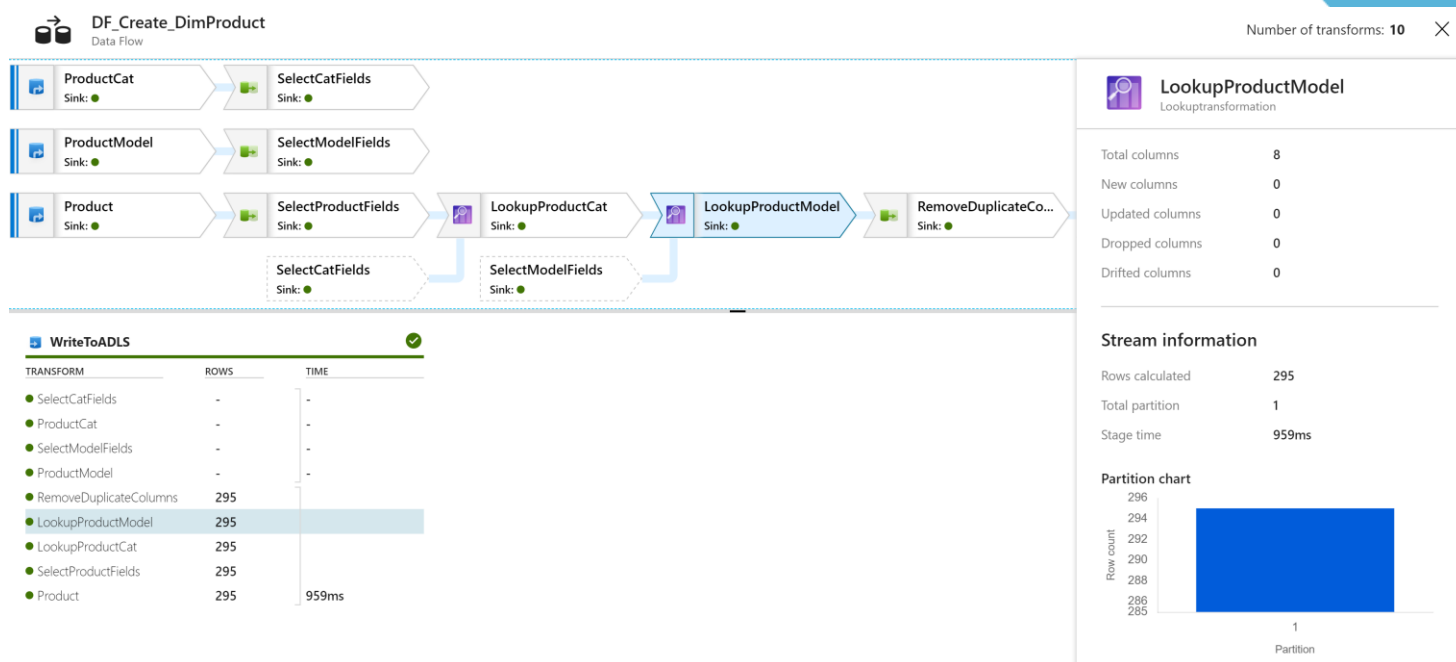
Here we can see how long each activity in our pipeline took to execute, which is useful if you have a long chain of transformations, but what we're interested in is the details behind the actual dataflow:

ACTIVITY NAME	ACTIVITY TYPE	ACTIONS	RUN START ↕	DURATION	STATUS
DF_Create_Di...	ExecuteDataFl...	  	06/27/2019 1:49:51 PM	00:06:18	 Succeeded

- Click the  to open the data flow results page, which breaks down each transformation and gives partitioning data as various spark transformations were made:



Lab 04 – Building a Mapping Data Flow



This page contains a wealth of information – highlight different transformations to see the number of rows processed, how long that stage lasted and how the data was partitioned. In this example, the data was very small and so we could perform everything on a single box. For larger examples, we can configure how datasets are distributed to optimise spark executor partitioning, which is very powerful indeed!

You've now got the basics for creating a Data Factory data flow, but there's a lot more to learn! Try out some of the other transformation types and, when you're ready, try using the DerivedColumn transformation to see the large number of functions available!